

Here we use cosine similarity to compare what annotations are similar.

We convert our tweets to a matrix which has for each column a word that appears in the corpus. The numbers in the matrix respond to the frequency of these words.

Using this matrix we can use cosine similarity which gives us how similar vectors are. Since our matrix is just a series of vectors, we can apply this to our matrix and in turn our tweets.

This section 2.3 ends with a heat map which tells us how similar the annotations are. If the colour is more white, that means those two annotations are similar.

In our heatmap, we can see annotation 1 and 13 are similar, with a cosine similarity score of 0.95. These annotations correspond to "calling out or correction" and "sarcasm".

First we take all the text for each group and put it into one row as we are asked to do this not on an individual tweet basis.

```
In [321]: dfg = df.groupby("Annotation")
```

```
In [322]: column_names = ["Annotation", "Text"]
df_groups = pd.DataFrame(columns = column_names)
annots = []
fulltexts = []
for x, y in dfg:
    #print(x)
    #print(y['text'].str.cat(sep=', '))
    annots.append(x)
    fulltexts.append(y['text'].str.cat(sep=', '))
```

```
In [323]: dfg2 = pd.DataFrame({"Annotation": annots, "Text": fulltexts})
```

```
In [324]: dfg2
```

	Annotation	Text
0	ambiguous or hard to classify	are hand dryers effective in killing the new c...
1	calling out or correction	no, lime juice won't immunize you: some people...
2	commercial activity or promotion	visit https://t.co/gqr287l0w for new jersey #...
3	conspiracy	do you think the coronavirus is a natural occu...
4	emergency	bhfp is in need of essential supplies due to c...
5	fake cure	"anti-vaxers have the cure to the coronavirus..."
6	fake treatment	i currently have the flu (haven't been tested ...
7	false fact or prevention	if you went to dr. ricciardi in east orange as...
8	false public health response	italy allows malaria and hiv drugs for coronav...
9	irrelevant	pages parody! plus, joe talks coronavirus and h...
10	news	health authorities have identified a new coron...
11	panic buying	as the demand for hand sanitisers soars in the...
12	politics	i hope your father recovers from the virus as ...
13	sarcasm or satire	i protected myself from the coronavirus by wea...
14	true prevention	wash your hands as often as possible #cureforc...
15	true public health response	delhi hc restrains publication of certain adve...

Then we convert the text into a matrix where each word is a column and it gets a value 1 if it appears in the annotation.

```
In [325]: dt2 = CountVectorizer().fit_transform(dfg2['Text'])
```

We use cosine similarity to produce what can be interpreted as akin to a correlation matrix.

```
In [326]: from sklearn.metrics.pairwise import cosine_similarity
sim1 = pd.DataFrame(cosine_similarity(dt2,dt2))
```

```
In [327]: sim1
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	1.000000	0.951407	0.710862	0.928218	0.762031	0.781331	0.755302	0.888132	0.374929	0.894340	0.764297	0.894340	0.764297	0.894340	0.764297	0.894340
1	0.951407	1.000000	0.768467	0.929891	0.811118	0.807141	0.771616	0.914836	0.447666	0.898867	0.809751	0.898867	0.809751	0.898867	0.809751	0.898867
2	0.710862	0.768467	1.000000	0.719620	0.725741	0.647337	0.653725	0.824105	0.573720	0.742712	0.811023	0.742712	0.811023	0.742712	0.811023	0.742712
3	0.928218	0.929891	0.719620	1.000000	0.757918	0.796624	0.731955	0.858089	0.415083	0.893473	0.834455	0.893473	0.834455	0.893473	0.834455	0.893473
4	0.762031	0.811118	0.725741	0.757918	1.000000	0.678575	0.672576	0.814522	0.436440	0.734060	0.733205	0.734060	0.733205	0.734060	0.733205	0.734060
5	0.781331	0.807141	0.647337	0.796624	0.678575	1.000000	0.785798	0.747627	0.463483	0.487822	0.735319	0.463483	0.487822	0.735319	0.463483	0.487822
6	0.755302	0.771616	0.653725	0.731955	0.672576	0.785798	1.000000	0.744831	0.435436	0.693030	0.693030	0.693030	0.693030	0.693030	0.693030	0.693030
7	0.888132	0.914836	0.824105	0.858089	0.814522	0.747627	0.744831	1.000000	0.463483	0.864897	0.804748	0.463483	0.864897	0.804748	0.864897	0.804748
8	0.374929	0.447666	0.573720	0.415083	0.436440	0.487822	0.435436	0.463483	1.000000	0.373890	0.574430	0.373890	0.574430	0.373890	0.574430	0.373890
9	0.894340	0.898867	0.742712	0.893473	0.734060	0.735319	0.693030	0.864897	0.373890	1.000000	0.813790	0.373890	1.000000	0.813790	0.373890	1.000000
10	0.764297	0.809751	0.811023	0.834455	0.733205	0.722016	0.693619	0.804748	0.574430	0.813790	1.000000	0.574430	0.813790	1.000000	0.574430	0.813790
11	0.821555	0.858587	0.781626	0.854849	0.733672	0.716538	0.682270	0.837533	0.450327	0.897840	0.862490	0.897840	0.862490	0.897840	0.862490	0.897840
12	0.894777	0.918822	0.739685	0.901658	0.778550	0.788127	0.737589	0.854271	0.408884	0.878635	0.800077	0.878635	0.800077	0.878635	0.800077	0.878635
13	0.935070	0.953431	0.710877	0.914212	0.789106	0.786461	0.725426	0.890421	0.376851	0.913416	0.762752	0.376851	0.913416	0.762752	0.376851	0.913416
14	0.845954	0.885512	0.832002	0.871755	0.819130	0.689336	0.684972	0.905747	0.442588	0.832903	0.769123	0.442588	0.832903	0.769123	0.442588	0.832903
15	0.819947	0.867512	0.838570	0.874116	0.813748	0.787059	0.741265	0.871630	0.549838	0.858732	0.920375	0.549838	0.858732	0.920375	0.549838	0.920375

We display this in a heat map, with lighter colours representing a higher similarity. For example, 13 and 1 have a value of 0.95. These groups of tweets are very similar. What groups are they?

```
In [328]: annots[1], annots[13]
```

```
Out[328]: ('calling out or correction', 'sarcasm or satire')
```

So calling out/correction and sarcasm tweets are similar.

```
In [329]: plt.figure(figsize=(12, 8))
plt.xlabel("group")
plt.ylabel("group")
sns.heatmap(sim1, annot=True)
```

```
Out[329]: <AxesSubplot>
```



Use the following as a key to link the numbers with annotations

```
In [330]: annots
```

```
Out[330]: ['ambiguous or hard to classify',
'calling out or correction',
'commercial activity or promotion',
'conspiracy',
'emergency',
'fake cure',
'fake treatment',
'false fact or prevention',
'false public health response',
'irrelevant',
'news',
'panic buying',
'politics',
'sarcasm or satire',
'true prevention',
'true public health response']
```

2.4 What is the overall distribution of nouns and noun phrases, and verbs and verb phrases, in the entire corpus?

Using the tables below it can be seen that the amount of nouns is 28,730 and nouns are the most common word type. We display these word types graphically as a percentage below.

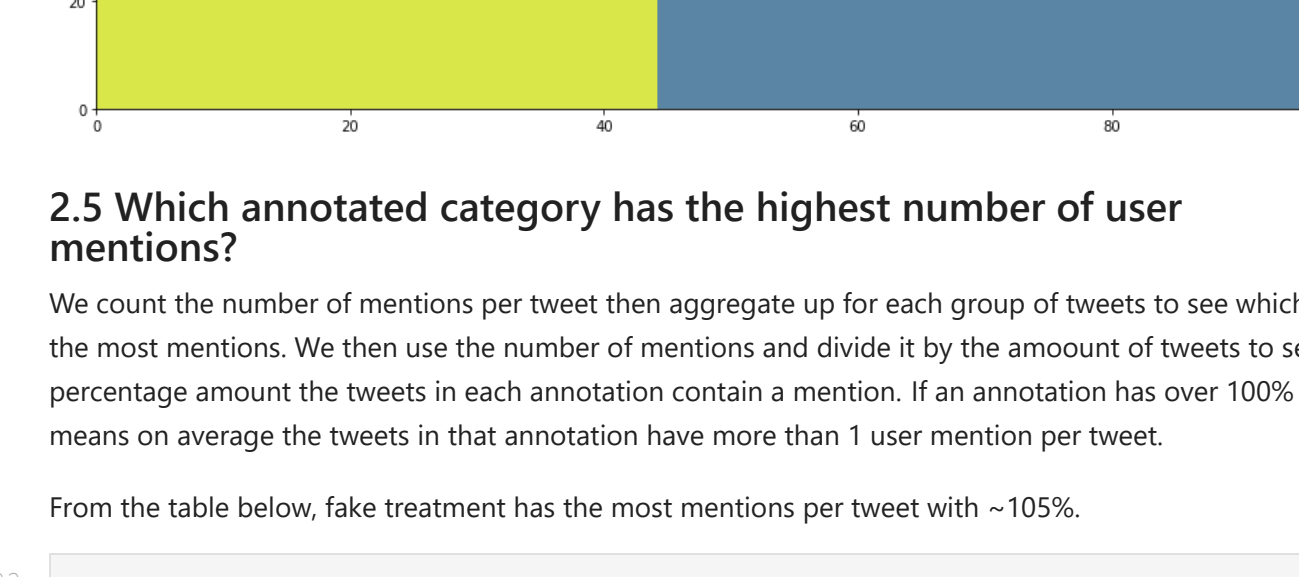
We sum up the number of each word type and display as a percentage of the entire dataset

```
In [331]: dfb3 = df1.sum()
dfb3[1:8]
```

```
Out[331]: adjs_num      20387
verbs_num      12927
nouns_num      28730
noun_phrases_num  5417
verb_phrases_num  1175
adj_noun_phrases_num  5606
entities_num     938
dtype: object
```

```
In [332]: #pip install squarify
import squarify
label= ["Adjectives", "Verbs", "Nouns", "Noun Phrases", "Verb Phrases", "AdjectiveNoun", "VerbNoun"]
perc = [(i/(dfb3[1:8].sum()*100),5,2f"%") for i in dfb3[1:8]]
lbl = ['%'] * len(perc)
fig = plt.figure(figsize=(16, 7))
squarify.plot(size=perc, label=lbl, alpha=0.8)
plt.axis('off')
plt.show()
```

```
Out[332]: <AxesSubplot>
```



2.5 Which annotated category has the highest number of user mentions?

We count the number of mentions per tweet then aggregate up for each group of tweets to see which has the most mentions. We then use the number of mentions and divide it by the amount of tweets to see the percentage amount the tweets in each annotation contain a mention. If an annotation has over 100% that means on average the tweets in that annotation have more than 1 user mention per tweet.

From the table below, fake treatment has the most mentions per tweet with ~105%.

```
In [333]: df['numuser'] = df.text.str.count("@user")
```

```
In [334]: dfb4 = df[df.columns[1p.concatenate([range(2,3),range(29,30)])]]
```

Sum up the user mentions in each annotation

```
In [335]: dfb4a = dfb4.groupby('annotation', as_index=False).sum()
dfb4a
```

```
Out[335]:
```

	annotation	numuser
0	ambiguous or hard to classify	38
1	calling out or correction	381
2	commercial activity or promotion	6
3	conspiracy	166
4	emergency	8
5	fake cure	71
6	fake treatment	22
7	false fact or prevention	51
8	false public health response	0
9	irrelevant	22
10	news	31
11	panic buying	17
12	politics	190
13	sarcasm or satire	78
14	true prevention	32
15	true public health response	53

Convert this to a percentage of the amount of tweets in each annotation

```
In [336]: totals = dfb4.groupby('annotation', as_index=False).agg('count')
dfb4a['percent'] = (dfb4a['numuser']/totals['numuser'])*100
```

From the table below, fake treatment has the most mentions per tweet with ~105%.

```
In [337]: dfb4a #one entry can have more than mention, hence fake treatment has more mentions
```

```
Out[337]:
```

	annotation	numuser	percent
0	ambiguous or hard to classify	38	33.628319
1	calling out or correction	381	31.936295
2	commercial activity or promotion	6	20.689655
3	conspiracy	166	32.549020
4	emergency	8	50.000000
5	fake cure	71	71.713913
6	fake treatment	22	104.761905
7	false fact or prevention	51	20.158103
8	false public health response	0	0.000000
9	irrelevant	22	20.183466
10	news	31	33.695652
11	panic buying	17	26.153846
12	politics	190	43.378955
13	sarcasm or satire	78	19.447355
14	true prevention	32	20.000000
15	true public health response	53	35.099338

2.5 Which annotated category uses the most hashtags in a tweet on average?

We use the same technique above but with hashtags. We can see Commercial activity has the most hashtags per tweet at ~24% (that's 24 hashtags per tweet)

```
In [338]: df['hash_count'] = df['hashtags_orig'].str.len()
```

```
In [339]: dfb5 = df[df.columns[1p.concatenate([range(2,3),range(30,31)])]]
dfb5a = dfb5.groupby('annotation', as_index=False).sum()
totals = dfb5.groupby('annotation', as_index=False).agg('count')
dfb5a['totals'] = (dfb5a['hash_count']/totals['hash_count'])*100
```

Table of percentage of hashtag in each annotation.

```
In [340]: dfb5a
```

```
Out[340]:
```

	annotation	hash_count	totals	percent
0	ambiguous or hard to classify	116	113	102.654867
1	calling out or correction	1017	1193	85.247276
2	commercial activity or promotion	70	291	24.1379310
3	conspiracy	739	510	144.901961
4	emergency	16	16	100.000000
5	fake cure	109	92	118.478261
6	fake treatment	16	21	76.190476
7	false fact or prevention	365	253	144.268775
8	false public health response	0	3	0.000000
9	irrelevant	184	109	168.807339
10	news	72	92	78.260870
11	panic buying	100	65	153.846154
12	politics	534	438	121.917808
13	sarcasm or satire	238	397	59.949622
14	true prevention	210	160	131.250000
15	true public health response	183	151	121.192053

2.6 What are the top 3 hashtags for each category, including and excluding the hashtags used to filter the tweets?

We print out tables below showing the top hashtags for each category including and excluding the hashtags. For example, for the annotation "fake cure" the most used hashtag was "coronil"

Create a program that counts the frequency of word appearance.

```
In [341]: from collections import Counter
def count_words(df, column='tokens', preprocess=None, min_freq=2):
    # process tokens and update counter
    def update(doc):
        tokens = doc if preprocess is None else preprocess(doc)
        counter.update(tokens)
    # create counter and run through all data
    counter = Counter()
    df[column].map(update)
    # transform counter into a DataFrame
    freq_df = pd.DataFrame.from_dict(counter, orient='index', columns=['freq'])
    freq_df = freq_df.query('freq >= fmin_freq')
    freq_df.index.name = 'token'
    return freq_df.sort_values('freq', ascending=False)
```

```
In [342]: hash_freq_tab = count_words(df, 'hashtags_orig').head(3)
#count_words(df, 'count_p').head(10).plot(kind='bar').invert_yaxis()
```

Top 3 hashtags per category including all tweets

```
In [343]: for x in dfb5a['annotation']:
    print(x)
    print(count_words(df.loc[df['annotation'] == x], 'hashtags_orig').head(3))
    print('\n')
```

ambiguous or hard to classify

```
token
covid19      17
coronavirus  14
bioweapon    4
```

calling out or correction

```
token
coronavirus  179
covid19      167
coronavirusfacts  23
```

commercial activity or promotion

```
token
coronavirus  111
covid19      63
bioweapon    62
```

conspiracy

```
token
coronavirus  111
covid19      63
bioweapon    62
```

emergency

```
token
coronavirus  4
covid19      2
donateplasmarecueallife  2
```

fake cure

```
token
coronil      8
covid19      6
coronavirus  6
```

fake treatment

```
token
hydroxychloroquine  2
coronavirus  2
```

false fact or prevention

```
token
coronavirus  39
covid19      35
covid19_19  18
```

false public health response

```
Empty DataFrame
Columns: [freq]
Index: []
```

irrelevant

```
token
coronavirus  29
covid19      21
cooking      4
```

news

```
token
covid19      3
coronavirus  18
covid      3
```

panic buying

```
token
coronavirus  20
covid19      7
covid19_19  4
```

politics

```
token
coronavirus  60
covid19      54
trump        37
```

sarcasm or satire

```
token
coronavirus  38
covid19      34
socialdistancing  6
```

true prevention

```
token
covid19      48
coronavirus  26
covid19_19  6
```

true public health response

```
token
hydroxychloroquine  6
covid19      28
coronavirus  23
who            4
```

3. Task C – Affect Analysis

In this analysis we use a dictionary known as Vader, which is a popular dictionary to use with twitter data.

We link the words in our tweets to that dictionary and it then gives each word a positive or negative score.

The word "good" in a tweet for example would get a score +1, and the word "horrible" might get a score of -2. Using this method, we can see for each annotation how positive/negative it is.

At the end of this section there is a plot which shows how positive or negative each annotation is. Scores above 0 are deemed positive with a maximum positive score of 1. Scores below 0 are deemed negative with a minimum negative score of -1.

```
In [345]: from nltk.sentiment.vader import SentimentIntensityAnalyzer
sid = SentimentIntensityAnalyzer()
```

```
In [346]: dfg2['Text'] ###here is the column we'll do the analysis on
```

```
Out[346]:
```

```
0 are hand dryers effective in killing the new c...
1 no, lime juice won't immunize you: some people...
2 visit https://t.co/gqr287l0w for new jersey #...
3 do you think the coronavirus is a natural occu...
4 bhfp is in need of essential supplies due to c...
5 "anti-vaxers have the cure to the coronavirus..."
6 i currently have the flu (haven't been tested ...
7 if you went to dr. ricciardi in east orange as...
8 italy allows malaria and hiv drugs for coronav...
9 pages parody! plus, joe talks coronavirus and h...
10 health authorities have identified a new coron...
11 as the demand for hand sanitisers soars in the...
12 i hope your father recovers from the virus as ...
13 i protected myself from the coronavirus by wea...
14 wash your hands as often as possible #cureforc...
15 delhi hc restrains publication of certain adve...
Name: Text, dtype: object
```

This prepares a dataset with the group and all the text within the group.

```
In [347]: column_names = ["Annotation", "Text"]
df_groups = pd.DataFrame(columns = column_names)
annots = []
fulltexts = []
for x, y in dfg
```


- Fake Cure
- Fake Treatment
- False Fact or Prevention

Information

Annotations:

- True treatment
- True prevention
- True public health response
- Calling out or correction
- Emergency
- Irrelevant
- Commercial activity or promotion
- Panic buying
- Politics

We exclude "sarcasm" from this analysis. Sarcasm could confuse the classifier as the classifier probably won't be able to pick up on a sarcastic tone and may interpret sarcasm as misinformation

We use these tweets to split into test and train.

The train tweets will be used to build a Support Vector Machine model (SVM models are good for sparse data which we see in text data, that is matrices with many 0s).

We will then use the test tweets to see how well the model can correctly label these tweets as information or mis-information.

Since we know these tweets are classed as information/mis-information, we can cross check what the model predicts with the reality.

In this model, we build a model with 82.4% accuracy.

We also carried out a grid search later on in this section. This takes many parameters that can be inputted into the model to see which gives the best accuracy score. Using this, we used the best parameters in our model.

```
In [383]_ ## merge tweets
df

misinfo = ["conspiracy",
           "fake cure",
           "fake treatment",
           "false fact or prevention"]
no_misinfo = ["true treatment",
              "true prevention",
              "true public health response",
              "calling out or correction",
              "emergency",
              "irrelevant",
              "commercial activity or promotion",
              "panic buying",
              "politics"]
```

Assign each tweet a misinfo or no misinfo label

```
In [384]_ # create a list of our conditions
conditions = [
    (df["annotation"].isin(misinfo)),
    (df["annotation"].isin(no_misinfo)),
    ~(df["annotation"].isin(misinfo)) & ~(df["annotation"].isin(no_misinfo))
]

# create a list of the values we want to assign for each condition
values = ['misinfo', 'no_misinfo', 'n/a']

# create a new column and use np.select to assign values to it using our lists as arg
df["class1"] = np.select(conditions, values)
```

A table showing an annotation and how it's classified (class1). Irrelevant is classified as no_misinfo which we want.

```
In [385]_ with pd.option_context('display.max_rows', None, 'display.max_columns', None):
    print(df[["annotation", 'class1']].head())

annotation    class1
0      irrelevant  no_misinfo
1      irrelevant  no_misinfo
2      irrelevant  no_misinfo
3      irrelevant  no_misinfo
4      politics    no_misinfo
```

Drop the tweets we don't deem relevant to the analysis.

```
In [386]_ #drop n/a
dfdl = df[df.class1 != 'n/a']
```

The column 'text9' is after our pre-processing steps. We change it to a list to a string so it can be used in the classifier.

```
In [387]_ #hello = ' '.join(dfdl['text9'])
dfdl['text9']
dfdl['text10'] = [' ' ' '.join(x) for x in dfdl['text9']]

<ipython-input-387-59c8dd610ef5b3:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/10min.html#setting-with-copy
dfdl['text10'] = [' ' ' '.join(x) for x in dfdl['text9']]
```

```
In [388]_ #dfdl['text10'][:352]  #= dfdl['text10'].map(remove_hash)
df['text8'][:352]
```

```
Out[388]_ [hey,
8user,
cocaine,
gay,
orgies,
similar,
diplomatic,
initiatives,
stop,
baya,
bitch,
balls,
imranhantti,
pakistan,
coronaelect,
coronabreak,
coronaviruspandemic,
coronavirius,
searc,
https://t.co/55n0bn9zrv]
```

Split the data into train and test. Train data is used to create the model. We then use the test data to see if the model is accurate on data we didn't use in the model to see if it can make predictions on new data.

```
In [389]_ #use text6
X_train, X_test, Y_train, Y_test = train_test_split(dfdl['text10'],
                                                    dfdl['class1'],
                                                    test_size=0.2,
                                                    random_state=42,
                                                    stratify=dfdl['class1']
                                                    )

print('Size of Training Data ', X_train.shape[0])
print('Size of Test Data ', X_test.shape[0])

Size of Training Data 2503
Size of Test Data 626
```

Transform the test data into a matrix of numbers, where each word gets a number of how frequently it appears in each class of info and misinfo. The grid search below, which tells us the best parameters to use for this model, tells us that an ngram range of 1,3 is the best. That means we include phrases of 1, 2 and 3 words long in the analysis. For example for a 3 gram phrase, "coronavirus is bad" would be included. For a 2 gram phrase "go away" would be included.

```
In [390]_ from sklearn.feature_extraction.text import TfidfVectorizer
tfidf = TfidfVectorizer(min_df = 10, ngram_range=(1,3), stop_words="english")
X_train_tf = tfidf.fit_transform(X_train)
```

We use a Support Vector Machine to classify as this is better with text data as it's sparse data (the matrix has a lot of 0s in, hence it is sparse).

```
In [391]_ #from sklearn.svm.LinearSVC import LinearSVC
from sklearn.svm import LinearSVC
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import make_classification
model = LinearSVC(random_state=0, tol=1e-5)
model.fit(X_train_tf, Y_train)
```

```
Out[391]_ LinearSVC(random_state=0, tol=1e-05)
```

```
In [392]_ LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=10000,
multi_class='ovr', penalty='l2', random_state=0, tol=1e-05,
verbose=0)
```

```
Out[392]_ LinearSVC(max_iter=10000, random_state=0, tol=1e-05)
```

We test the model on our test data to see how accurate it is. For each tweet we let the model predict if it's information or misinformation and we compare it to the actual label. What percentage of the time is it correct, the accuracy score tells us that, 82.4%

```
In [393]_ X_test_tf = tfidf.transform(X_test)

Y_pred = model.predict(X_test_tf)
print('Accuracy Score - ', accuracy_score(Y_test, Y_pred))

Accuracy Score - 0.8242811501597445
```

This is a confusion matrix which tells us how well the model performs.

For the test data, and for the misinformation tweets, the model gets it right 104 times, but gets it wrong 71 times. That's not so good. That means given the tweet has misinformation, it correctly identifies it as misinformation 59% of the time.

For the no_misinformation, it gets it wrong 39 times but gets it right 412 times. That's good. That means if a tweet comes in that has no information, we can correctly determine using the model with a rate of 91% that it has no information.

So it's good at detecting tweets with no misinformation but not as good at picking up tweets with misinformation.

```
In [394]_ from sklearn.metrics import plot_confusion_matrix
plot_confusion_matrix(model,X_test_tf,
                      Y_test, values_format='d',
                      cmap=plt.cm.Blues)

plt.show()
```



Here are some examples where the model gets it right. You can see the columns actual and predicted match up.

```
In [395]_ # Create a DataFrame combining the Title and Description,
# Actual and Predicted values that we can explore
frame = {'text': X_test, 'actual': Y_test, 'predicted': Y_pred }
result = pd.DataFrame(frame)

result[(result['actual'] == 'no_misinfo') | (result['actual'] == 'misinfo')] & (result['actual'] == result['predicted'])
```

```
Out[395]_      text    actual  predicted
115  bacterial horror hot air hand dryers https://t...  no_misinfo  no_misinfo
607  alcohol disinfectant chlorine effective clean...  no_misinfo  no_misinfo
2149  relative maga actually covid 19 fake news trum...  no_misinfo  no_misinfo
1338  day trump give suggestion case people intima...  no_misinfo  no_misinfo
```

Here are some examples where the model gets it wrong. The table shows the actual label, which is what the tweet actually is, and the predicted label, what the model predicts the tweet to be.

```
In [396]_ result[(result['actual'] == 'misinfo') & (result['actual'] != result['predicted'])].\
frame[['text', 'actual', 'predicted']]

Out[396]_      text    actual  predicted
1249  @user see precious tweet family corona virus...  misinfo  no_misinfo
3605  look recent study claim people hospitalize w...  misinfo  no_misinfo
2159  comment researcher find cheap widely availab...  misinfo  no_misinfo
2958  big concern pay attention fact covid 19 airbor...  misinfo  no_misinfo
```

Here is an example tweet that the classifier gets wrong

```
In [399]_ dfdl['text'][:3605]
```

```
Out[399]_ 'I looked at the recent study that claims that people hospitalized w/covid-19 did better on hydroxy than without.\n\nI people given hydroxy were also more likely to be given steroids (know to be helpful). 79% vs 36%\n2. people given hydroxy were on average 5 years younger.\n\nI/1/2* https://t.co/6957bazin8'
```

4.1 Thoughts

Although the model performs reasonably well, how it classifies tweets with misinformation isn't so accurate.

Including "conspiracy" annotation is debatable in the misinformation, since it doesn't really give advice on how to deal with covid. The model was run without the conspiracy annotation and unfortunately the number of misinformation tweets was too low to build a stable model. If this analysis were to be run again more tweets about misinformation would make a big improvement.

4.2 Appendix

This is a grid search which tells us the best parameters to use. This is commented out because it only needs to be run once and the parameters from this and used in the SVM model creation in section 4.

```
In [398]_ # from sklearn.pipeline import Pipeline
# from sklearn.model_selection import GridSearchCV
# training_pipeline = Pipeline(
#     steps=[('tfidf', TfidfVectorizer(stop_words="english")),
#            ('model', LinearSVC(random_state=42, tol=1e-5))]
# )
# grid_param = [
#     {'tfidf_min_df': [5, 10],
#      'tfidf_ngram_range': [(1, 3), (1, 6)],
#      'model_penalty': ['l2'],
#      'model_loss': ['hinge'],
#      'model_max_iter': [10000]
#     }, {
#         'tfidf_min_df': [5, 10],
#         'tfidf_ngram_range': [(1, 3), (1, 6)],
#         'model_c': [1, 10],
#         'model_tol': [1e-2, 1e-3]
#     }
# ]
# grid_search_processor = GridSearchCV(estimator=training_pipeline,
#                                     param_grid=grid_param,
#                                     cv=5)
# grid_search_processor.fit(dfdl['text10'], dfdl['class1'])

# best_params = grid_search_processor.best_params_
# print('Best alpha parameter identified by grid search ', best_params)

# best_result = grid_search_processor.best_score_
# print('Best result identified by grid search ', best_result)
```